

# MPI IMPLEMENTATION OF K-MEANS CLUSTERING

---

Sepideh Hayati  
Alireza Khornegah

Spring 2025



# WHAT IS K-MEANS CLUSTERING?



K-means is an unsupervised learning method for clustering data points.



**Defenition**




**How doese it work?**



**Its Goal**



# DATASETS GENERATION



## How to Generate Datasets

In Python using NumPy



## Feature of Datasets

Three cluster centers  
With a standard deviation of 400  
2D random points



## Size of Datasets

20K, 40K, 80K


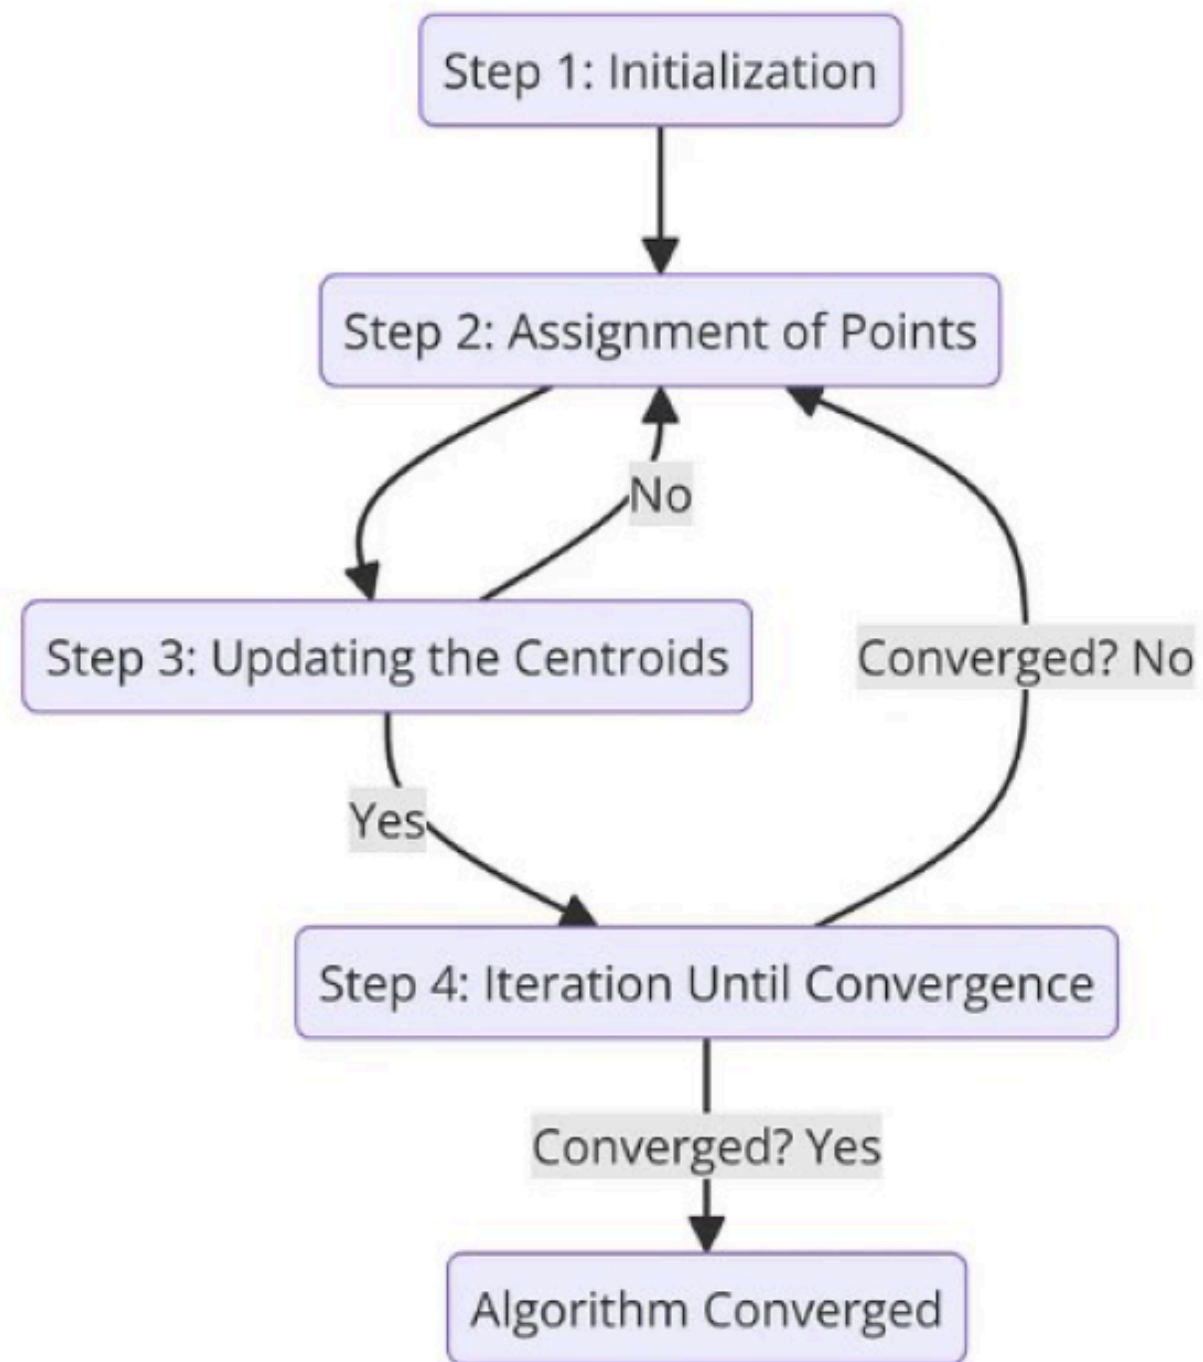
# SERIAL IMPLEMENTATION

---

# Algorithm Steps



Given a dataset and the parameters below, the goal is to assign each point to one of K clusters and compute the final centroids.



# OBSERVATION FOR PARALLELIZATION

- **Data Distribution and Broadcasting**
- **Initial Centroid Selection and Distribution**
- **Workload Partitioning, Distance Calculation and Cluster Assignment**
- **Parallel Centroid Update with Reduction**

# Data Distribution and Broadcasting

## Parallel Version

- Only one processor reads the dataset from the file shares the data with all the others using MPI.
- MPI Functions Used: MPI Bcast

## Serial Version

- Entire data is read by one process, no sharing or distribution needed.

# Initial Centroid Selection and Distribution

## Parallel Version

- Root process selects the initial centroids (starting points for clusters) and shares them.
- MPI Functions Used: MPI Bcast

## Serial Version

- Initial centroids are selected and used by one process, without sharing.

# Workload Partitioning, Distance Calculation and Cluster Assignment

## Parallel Version

The workload is divided among all processes:

- Each process handles a subset of the data points.
- Every process calculates distances only for its own chunk.
- Each point is assigned to the nearest cluster locally.
- MPI Functions Used: MPI Comm rank, MPI Comm size, MPI Allgatherv

## Parallel Centroid Update with Reduction

### Parallel Version

- Each process works on its own portion of the dataset
- Computes the local sum of coordinates for each cluster.
- Computes the local count of points assigned to each cluster.
- MPI Functions Used: MPI Allreduce

### Serial Version

- It loops through the entire dataset.
- For each point, it calculates the distance to all centroids.
- It then assigns that point to the closest cluster.
- stores the result in one process—no need for communication or data distribution.

### Serial Version

In the serial version, one process:

- Adds up all data points assigned to each cluster.
- Calculates the mean (average) to find the new centroid.




# AMDAHL'S LAW

To guess how much faster it can get, Amdahl's Law is used:

$$Speedup(N) = \frac{1}{S + \frac{P}{N}}$$

Sum of Time

Speedup: Actual and Amdal's law



There is a GAP!

This highlights the influence of factors such as communication overhead, memory bandwidth limitations, and hardware inefficiencies.

# PERFORMANCE ANALYSIS

---

# Machines

To measure the performance of the MPI-based solution, two types of clusters were set up on GCP:

## Fat Clusters:

- Two nodes
- Each with 4 vCPUs
- Intra-regional, both in the same region
- Location: us-central1-a and us-central1-b

## Light Clusters:


- Four nodes
- Each with 2 vCPUs
- Intra-regional, all in the same region
- Location: us-central1-a

# Fat Clusters:

## Strong scalability

Sum of Speed Up

Speedup by Number of Cores

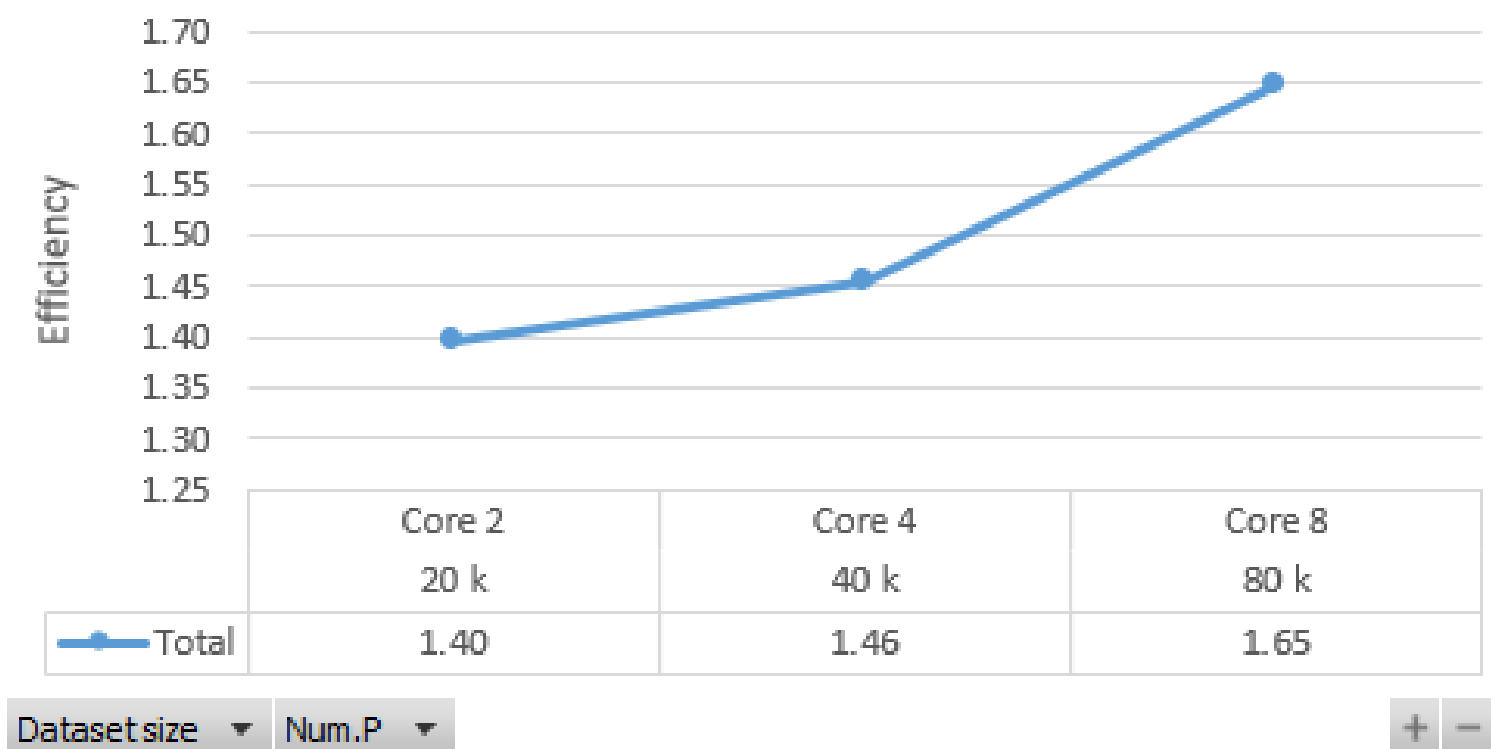


- By increasing the number of cores, the speedup improves.
- The speedup increases more slowly for larger datasets like 40K and 80K.
- This is likely due to MPI communication overhead, which grows as more cores are added, reducing the benefits of parallelism for larger datasets.

## Weak scalability

Sum of Efficiency


Efficiency



- As we increase both the dataset size and the number of cores proportionally, efficiency improves.
- This shows that the system handles larger workloads well when more resources are added, indicating good weak scalability on fat clusters.

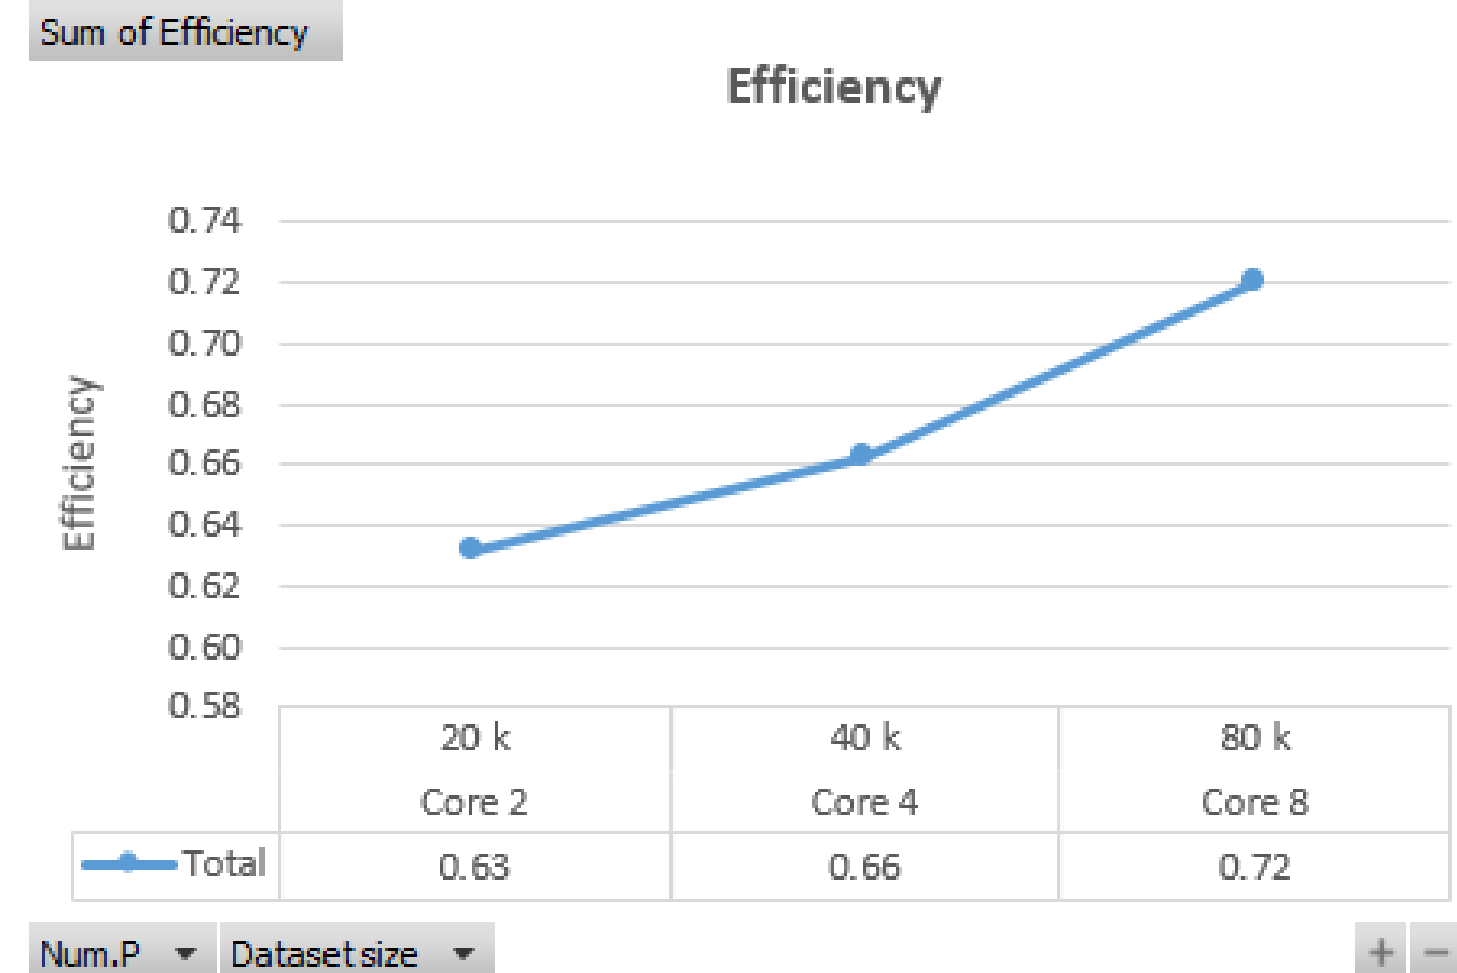
# Light Clusters:

## Strong scalability



- As the number of cores increases from 1 to 8, the speedup remains low and even drops slightly for all dataset sizes.
- This indicates poor strong scalability, due to the limited processing power of light clusters or Increased communication overhead

## Weak scalability



- Efficiency improves steadily as both dataset size and number of cores increase proportionally .
- The system handles larger workloads more efficiently when more resources are available.
- The system shows moderate weak scalability